

ormance-screeningassessmentcohort3

June 30, 2023

#Libraries

```
[1]: import pandas as pd #importing panda as pd
from scipy.stats import ttest_ind # for t-test
import matplotlib.pyplot as plt # for plotting (like bar plot,
    ↪ histogram, scatter plot)
import numpy as np
```

#Task1 Download the dataset and create an IPython Notebook or a Google Colab notebook for this assignment.

```
[2]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

#Task2 Load the dataset into a pandas dataframe and display the first 10 rows of the dataframe to gain familiarity with the data structure.

```
[3]: df = pd.read_csv('/content/drive/MyDrive/StudentPerformance Dataset/
    ↪ StudentsPerformance.csv')
result = df.head(10) #This function will return first 10 rows
print(result)
```

	gender	race/ethnicity	parental level of education	lunch	\
0	female	group B	bachelor's degree	standard	
1	female	group C	some college	standard	
2	female	group B	master's degree	standard	
3	male	group A	associate's degree	free/reduced	
4	male	group C	some college	standard	
5	female	group B	associate's degree	standard	
6	female	group B	some college	standard	
7	male	group B	some college	free/reduced	
8	male	group D	high school	free/reduced	
9	female	group B	high school	free/reduced	

	test preparation course	math score	reading score	writing score
0	none	72	72	74
1	completed	69	90	88

2	none	90	95	93
3	none	47	57	44
4	none	76	78	75
5	none	71	83	78
6	completed	88	95	92
7	none	40	43	39
8	completed	64	64	67
9	none	38	60	50

```
[4]: print(df)
```

	gender	race/ethnicity	parental level of education	lunch	\
0	female	group B	bachelor's degree	standard	
1	female	group C	some college	standard	
2	female	group B	master's degree	standard	
3	male	group A	associate's degree	free/reduced	
4	male	group C	some college	standard	
..	
995	female	group E	master's degree	standard	
996	male	group C	high school	free/reduced	
997	female	group C	high school	free/reduced	
998	female	group D	some college	standard	
999	female	group D	some college	free/reduced	

	test preparation course	math score	reading score	writing score
0	none	72	72	74
1	completed	69	90	88
2	none	90	95	93
3	none	47	57	44
4	none	76	78	75
..
995	completed	88	99	95
996	none	62	55	55
997	completed	59	71	65
998	completed	68	78	77
999	none	77	86	86

[1000 rows x 8 columns]

#Task3 Perform exploratory data analysis:

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                1000 non-null   object
```

```

1  race/ethnicity          1000 non-null  object
2  parental level of education 1000 non-null  object
3  lunch                   1000 non-null  object
4  test preparation course   1000 non-null  object
5  math score               1000 non-null  int64
6  reading score            1000 non-null  int64
7  writing score             1000 non-null  int64
dtypes: int64(3), object(5)
memory usage: 62.6+ KB

```

df.info tells us that there are 1000 rows, with no row to have any null values. And it has total 8 columns. Also there we can find data types of each column.

```
[6]: df.dtypes # for specifically checking data types
```

```

[6]: gender                object
     race/ethnicity         object
     parental level of education object
     lunch                  object
     test preparation course  object
     math score              int64
     reading score           int64
     writing score            int64
     dtype: object

```

```
[7]: df.shape # for specifically checking rows and columns
```

```
[7]: (1000, 8)
```

```
[8]: df.isnull() #checking if there is any null value
```

```

[8]:      gender  race/ethnicity  parental level of education  lunch  \
0      False                False                False  False
1      False                False                False  False
2      False                False                False  False
3      False                False                False  False
4      False                False                False  False
..      ...                  ...                  ...    ...
995    False                False                False  False
996    False                False                False  False
997    False                False                False  False
998    False                False                False  False
999    False                False                False  False

      test preparation course  math score  reading score  writing score
0                False        False        False        False
1                False        False        False        False
2                False        False        False        False

```

3	False	False	False	False
4	False	False	False	False
..
995	False	False	False	False
996	False	False	False	False
997	False	False	False	False
998	False	False	False	False
999	False	False	False	False

[1000 rows x 8 columns]

```
[9]: df.isnull().sum()
```

```
[9]: gender                0
     race/ethnicity        0
     parental level of education  0
     lunch                  0
     test preparation course  0
     math score             0
     reading score          0
     writing score           0
     dtype: int64
```

```
[10]: df.duplicated() #checking if there is ant duplicated value
```

```
[10]: 0      False
      1      False
      2      False
      3      False
      4      False
      ...
      995    False
      996    False
      997    False
      998    False
      999    False
      Length: 1000, dtype: bool
```

```
[11]: df.duplicated().sum()
```

```
[11]: 0
```

```
[12]: #.value_counts() function return a Series containing counts of unique values

print(df["gender"].value_counts())
print(df["race/ethnicity"].value_counts())
print(df["parental level of education"].value_counts())
```

```

print(df["lunch"].value_counts())
print(df["test preparation course"].value_counts())
print(df["math score"].value_counts())
print(df["reading score"].value_counts())
print(df["writing score"].value_counts())

```

```

female      518
male        482
Name: gender, dtype: int64
group C      319
group D      262
group B      190
group E      140
group A       89
Name: race/ethnicity, dtype: int64
some college      226
associate's degree  222
high school       196
some high school   179
bachelor's degree  118
master's degree     59
Name: parental level of education, dtype: int64
standard        645
free/reduced     355
Name: lunch, dtype: int64
none            642
completed       358
Name: test preparation course, dtype: int64
65      36
62      35
69      32
59      32
61      27
..
24       1
28       1
33       1
18       1
8         1
Name: math score, Length: 81, dtype: int64
72      34
74      33
64      32
67      30
73      30
..
28       1

```

```

26      1
17      1
32      1
40      1
Name: reading score, Length: 72, dtype: int64
74     35
70     33
68     31
73     28
80     27
..
28      1
35      1
19      1
22      1
23      1

```

Name: writing score, Length: 77, dtype: int64

.to_frame() function converts the series object to a dataframe

```
[13]: (df["gender"].value_counts()).to_frame()
```

```
[13]:
      gender
female    518
male     482
```

```
[14]: (df["gender"].value_counts()).to_frame().values
```

```
[14]: array([[518],
            [482]])
```

```
[15]: (df["race/ethnicity"].value_counts()).to_frame()
```

```
[15]:
      race/ethnicity
group C           319
group D           262
group B           190
group E           140
group A            89
```

```
[16]: (df["parental level of education"].value_counts()).to_frame()
```

```
[16]:
      parental level of education
some college                    226
associate's degree              222
high school                     196
some high school                179
bachelor's degree              118
```

```
[17]: (df["lunch"].value_counts()).to_frame()
```

```
[17]:          lunch
standard      645
free/reduced  355
```

```
[18]: (df["test preparation course"].value_counts()).to_frame()
```

```
[18]:          test preparation course
none                                642
completed                          358
```

```
[19]: (df["math score"].value_counts()).to_frame()
```

```
[19]:          math score
65             36
62             35
69             32
59             32
61             27
..           ...
24             1
28             1
33             1
18             1
8              1

[81 rows x 1 columns]
```

```
[20]: (df["reading score"].value_counts()).to_frame()
```

```
[20]:          reading score
72             34
74             33
64             32
67             30
73             30
..           ...
28             1
26             1
17             1
32             1
40             1

[72 rows x 1 columns]
```

```
[21]: (df["writing score"].value_counts()).to_frame()
```

```
[21]:      writing score
74          35
70          33
68          31
73          28
80          27
..          ...
28           1
35           1
19           1
22           1
23           1

[77 rows x 1 columns]
```

```
[22]: df.describe() #it will return the description of numeric data in a dataframe by default
```

```
[22]:      math score  reading score  writing score
count  1000.00000    1000.000000    1000.000000
mean    66.08900     69.169000     68.054000
std     15.16308     14.600192     15.195657
min      0.00000     17.000000     10.000000
25%     57.00000     59.000000     57.750000
50%     66.00000     70.000000     69.000000
75%     77.00000     79.000000     79.000000
max     100.00000    100.000000    100.000000
```

count returns non-empty values, our data has 1000 data rows and returns 1000-which means there is no empty values. Also, here max value and mean value has not much difference

```
[23]: df.corr()
```

<ipython-input-23-2f6f6606aa2c>:1: FutureWarning:

The default value of `numeric_only` in `DataFrame.corr` is deprecated. In a future version, it will default to `False`. Select only valid columns or specify the value of `numeric_only` to silence this warning.

```
[23]:      math score  reading score  writing score
math score      1.000000      0.817580      0.802642
reading score    0.817580      1.000000      0.954598
writing score    0.802642      0.954598      1.000000
```

`.corr()` method shows the correlation of the columns. Each column here is correlated with one

another in a good way, as all the values are closer to 1. It means that when one value increases, the other also goes up.

```
[24]: df.describe(include = ['object']) #this will return description of object data
      ↪types
```

```
[24]:      gender race/ethnicity parental level of education      lunch \
count      1000           1000           1000      1000
unique         2             5             6         2
top    female      group C      some college  standard
freq       518          319          226       645

      test preparation course
count           1000
unique            2
top             none
freq           642
```

here, count shows the non-null rows. unique shows the different values in each columns, like gender has two types- male and female. top shows the most common values in a column and freq shows the frequency of most common value.

```
[25]: def find_outliers_IQR(df):

      q1=df.quantile(0.25)

      q3=df.quantile(0.75)

      IQR=q3-q1

      outliers = df[((df<(q1-1.5*IQR)) | (df>(q3+1.5*IQR)))]

      return outliers
```

```
[26]: outliers = find_outliers_IQR(df['math score'])

print("number of outliers: "+ str(len(outliers)))

print("max outlier value: "+ str(outliers.max()))

print("min outlier value: "+ str(outliers.min()))

outliers
```

```
"number of outliers: "8
"max outlier value: "26
"min outlier value: "0
```

```
[26]: 17      18
      59      0
      145     22
      338     24
      466     26
      787     19
      842     23
      980      8
      Name: math score, dtype: int64
```

```
[27]: outliers = find_outliers_IQR(df['writing score'])

      print("number of outliers: "+ str(len(outliers)))

      print("max outlier value: "+ str(outliers.max()))

      print("min outlier value: "+ str(outliers.min()))

      outliers
```

```
"number of outliers: "5
"max outlier value: "23
"min outlier value: "10
```

```
[27]: 59      10
      76      22
      327     19
      596     15
      980     23
      Name: writing score, dtype: int64
```

```
[28]: outliers = find_outliers_IQR(df['reading score'])

      print("number of outliers: "+ str(len(outliers)))

      print("max outlier value: "+ str(outliers.max()))

      print("min outlier value: "+ str(outliers.min()))

      outliers
```

```
"number of outliers: "6
"max outlier value: "28
"min outlier value: "17
```

```
[28]: 59      17
      76      26
```

```
211    28
327    23
596    24
980    24
Name: reading score, dtype: int64
```

There are very little outliers present in the numeric values and I think they show the variations in the population, for this reason I kept the outliers as it is.

#Task4

```
[29]: #Create a bar plot to show the distribution of male and female students in the
      ↪dataset.
students = df['gender'].value_counts()
students
```


```
[29]: female    518
      male     482
      Name: gender, dtype: int64
```

```
[30]: x= students.index
      y= students.values
      plt.bar(x,y,color=('c','y'))

      #function for add title
      plt.title("Bar plot to show the distribution of male and female students in the
      ↪dataset.")
      plt.grid(False)
      #function to add label
      plt.xlabel("gender")
      plt.ylabel("count")

      # function to show the plot
      plt.show()
```

Bar plot to show the distribution of male and female students in the dataset.



we can see from the bar plot that most of the students in the dataset are of the gender female.

```
[31]: #Create a histogram to show the distribution of the scores in math, reading, and writing.
m_score = df[['math score']]
r_score = df[['reading score']] #here we are extracting the score columns and saving the dataframes in the variables
w_score = df[['writing score']]


# plotting first histogram
plt.hist(m_score, bins = 20, label='Math Score', alpha=.7, color='blue', edgecolor='black')

# plotting second histogram
plt.hist(r_score, bins = 20, label="Reading Score", color='pink', alpha=.7, edgecolor='black')

# plotting third histogram
plt.hist(w_score, bins = 20, label="Writing Score", color='green', alpha=.7, edgecolor='black')
plt.legend()#

# Showing the plot using plt.show()
```

```
plt.show()
```



in this histogram, we can see that there is overlap of the math, reading and writing score. The three columns has the values almost in the same range, but the 'math score' column has some lowest values, other than the other two columns. "reading score" and "writing score" columns also has some values far from the most frequent values. The most frequent score among the variables is between 65-75 which is in "math score" column. The second most frequent marks between 70-75, which is in "writing score" column. The highest score one got is in "writing score"

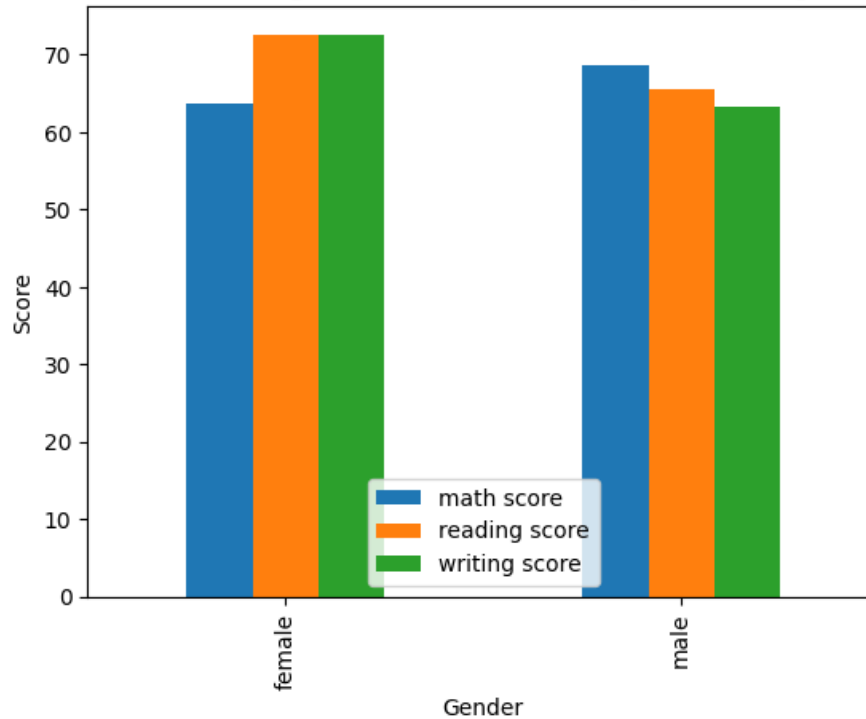
```
[32]: #Create a bar plot to show the average scores in math, reading, and writing for
      ↪ each gender.
avgScore = df.pivot_table(values=['math score', 'reading score', 'writing
      ↪ score'], index=['gender'], aggfunc='mean')
#creating dataframe using pivot_table of 3 columns and have set index = 'gender'
      ↪ because we want to group by female and male in pivot table
```

```
[33]: plt.figure(figsize=(10, 5))#specifying figure size
avgScore.plot(kind='bar')#plot() makes plots of dataframe avgScore
plt.title('Bar plot to show the average scores in math, reading, and writing
      ↪ for each gender')
plt.xlabel('Gender')
plt.ylabel('Score')
```

```
plt.legend( loc ="lower center")# is used to Place a legend on the axes.
# Showing the plot using plt.show()
plt.show()
```

<Figure size 1000x500 with 0 Axes>

Bar plot to show the average scores in math, reading, and writing for each gender



It is visible from the bar plot that female students have scored better than male students. Also, male students have scored better in maths than the female students. Male students have earned lowest scores in writing part.

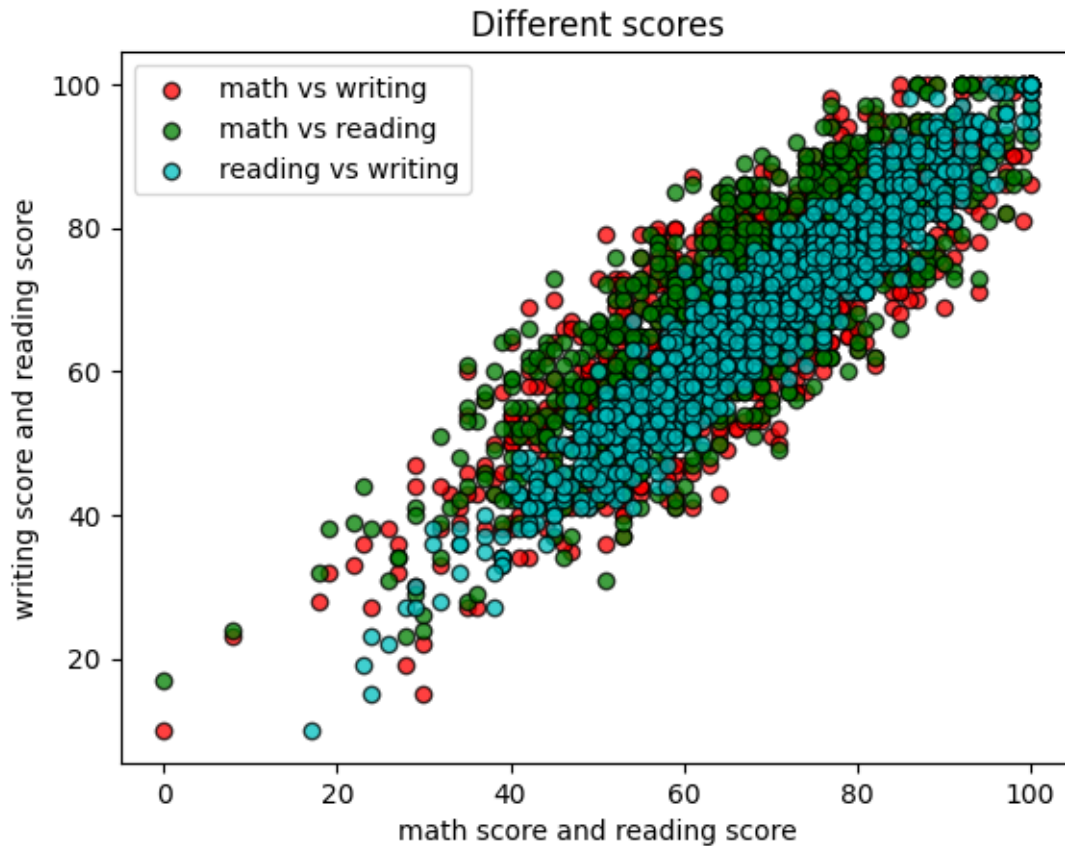
```
[78]: math_score = df['math score']
reading_score = df['reading score'] #saving the dataframes in the variables
writing_score = df['writing score']

#calling the scatter() method for showing the scatter plot.
plt.scatter(math_score, writing_score, color='r', edgecolor = 'black',
            linewidth = 1, alpha = 0.75, label = 'math vs writing')
plt.scatter(math_score, reading_score, color = 'g', edgecolor = 'black',
            linewidth = 1, alpha = 0.75, label = 'math vs reading')
plt.scatter(reading_score, writing_score, color = 'c', edgecolor = 'black',
            linewidth = 1, alpha = 0.75, label = 'reading vs writing')

plt.legend()# is used to Place a legend on the axes.
```

```
plt.title('Different scores')
plt.xlabel('math score and reading score')#specifying labels of x-axis
plt.ylabel('writing score and reading score')#specifying labels of y-axis
```

```
[78]: Text(0, 0.5, 'writing score and reading score')
```



The scatter plot shows moderately strong and positive relationships among different scores.

#Task5 Perform basic statistical analysis

```
[35]: #Calculate mean, median, and mode for math, reading, and writing scores.
mScore_mean = df['math score'].mean()
mScore_median = df['math score'].median()
mScore_mode = df['math score'].mode()
print('Mean: ', mScore_mean)
print('Median: ', mScore_mean)
print('Mode: ', mScore_mean)
```

```
Mean: 66.089
Median: 66.089
Mode: 66.089
```

```
[36]: rScore_mean = df['reading score'].mean()
rScore_median = df['reading score'].median()
rScore_mode = df['reading score'].mode()
print('Mean: ', rScore_mean)
print('Median: ', rScore_mean)
print('Mode: ', rScore_mean)
```

```
Mean: 69.169
Median: 69.169
Mode: 69.169
```

```
[37]: wScore_mean = df['writing score'].mean()
wScore_median = df['writing score'].median()
wScore_mode = df['writing score'].mode()
print('Mean: ', rScore_mean)
print('Median: ', rScore_mean)
print('Mode: ', rScore_mean)
```

```
Mean: 69.169
Median: 69.169
Mode: 69.169
```

as the mean, meadian and mode are all the same for each column, we can say that the distribution is symmetric

```
[38]: #Calculate the standard deviation for math, reading, and writing scores.
mScore_std = df['math score'].std()
rScore_std = df['reading score'].std()
wScore_std = df['writing score'].std()
print("Math Score Standard Deviation: ", mScore_std)
print("Reading Score Standard Deviation: ", rScore_std)
print("Writing Score Standard Deviation: ", wScore_std)
```

```
Math Score Standard Deviation: 15.16308009600945
Reading Score Standard Deviation: 14.600191937252216
Writing Score Standard Deviation: 15.19565701086965
```

```
[44]: #Test if there's a significant difference between male and female students'
#performance in math, reading, and writing using an appropriate statistical
↳test (e.g., t-test).
group1 = df[df['gender']=='female']
group2= df[df['gender']=='male']
```

```
<ipython-input-44-4a64b37c4fe8>:6: FutureWarning:
```

The default value of numeric_only in DataFrame.mean is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to

silence this warning.

```
[44]: math score      68.728216
      reading score   65.473029
      writing score    63.311203
      dtype: float64
```

```
[48]: group3 = df[df['gender']=='male'].mean() # to check the mean values of scores
      ↪where gender = male
      group3
```

<ipython-input-48-f6f74b357992>:1: FutureWarning:

The default value of `numeric_only` in `DataFrame.mean` is deprecated. In a future version, it will default to `False`. In addition, specifying `'numeric_only=None'` is deprecated. Select only valid columns or specify the value of `numeric_only` to silence this warning.

```
[48]: math score      68.728216
      reading score   65.473029
      writing score    63.311203
      dtype: float64
```

```
[47]: group4 = df[df['gender']=='female'].mean() # to check the mean values of scores
      ↪where gender = female
      group4
```

<ipython-input-47-f18e206aa25e>:1: FutureWarning:

The default value of `numeric_only` in `DataFrame.mean` is deprecated. In a future version, it will default to `False`. In addition, specifying `'numeric_only=None'` is deprecated. Select only valid columns or specify the value of `numeric_only` to silence this warning.

```
[47]: math score      63.633205
      reading score   72.608108
      writing score    72.467181
      dtype: float64
```

```
[41]: ttest_ind(group1['math score'], group2['math score'])
```

```
[41]: Ttest_indResult(statistic=-5.383245869828983, pvalue=9.120185549328822e-08)
```

```
[42]: ttest_ind(group1['reading score'], group2['reading score'])
```

```
[42]: Ttest_indResult(statistic=7.959308005187657, pvalue=4.680538743933289e-15)
```

```
[43]: ttest_ind(group1['writing score'], group2['writing score'])
```

```
[43]: Ttest_indResult(statistic=9.979557910004507, pvalue=2.019877706867934e-22)
```

as all the pValues are much less than 0.05, we can reject the null hypothesis and say that there is significant difference between the scores of male and female students.

In conclusion- the dataset has no null or duplicate values, math score, writing score and reading scores are correlated with one another, the dataset has more female students than male, female students' average score is greater than that of male students', but male students has done better in math than female students.